

# EE 622 Advanced Machine Learning

## Assignment 2 : Classification of cat vs. dog images

**Name: Immidisetti Rakhil**

**Roll No: 130102026**

**Department: ECE**

### Dependencies

Numpy, Scipy, Theano, Keras.

### Running Instructions

- Command for training the model:  
`THEANO_FLAGS=device=gpu,floatX=float32,lib.cnmem=0.8,mode=FAST_RUN`  
`python final_train.py`
- The data is shuffled and stored in numpy arrays in '.npz' format. So no need to normalize it to float everytime one runs the code.
- For every epoch, the weights are saved as '*params.h5*' which overwrites that of the previous epoch.
- The predictions for test data are in '*predictions.txt*' file.

### Report



#### Final architecture:

- The given data is shuffled and is split into training(0.8) and validation(0.2).
- Fits the model on data generated batch-by-batch by a Python generator. The generator is run in parallel to the model, for efficiency. This allows real-time data augmentation on images on CPU in parallel to training your model on GPU.
- The max generator queue size is 10. After every epoch the weights are saved and the metrics like loss and accuracy of that respective epoch are saved into a text file using callback.
- Used Image Data Generator in keras for data augmentation. Random horizontal flipping, shifting and rotation were done on the training images.
- Adadelat loss optimizer with a default learning rate of 0.1 and categorical cross-entropy as loss function were used.
- Activation function- LeakyRelu with alpha=0.3.



- A 8-layer architecture with 6 conv and 2 fc layers. 2 max-pooling layers were used after the first two conv layers.
- The output to the first fully connected layer is of shape 6x6x256.
- A high dropout of 0.7 was used in the final fully connected layers.

### Training methods and architectures:

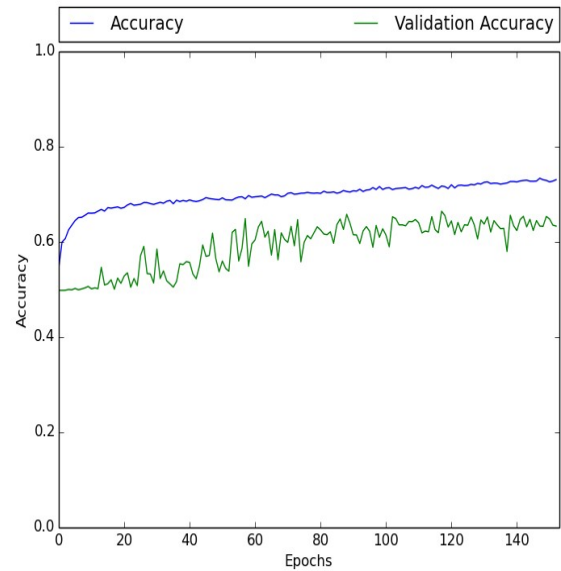
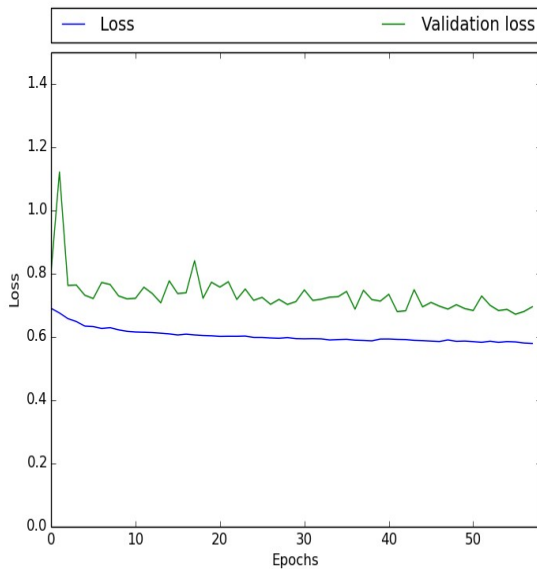
- First, I started out with a 7 layer architecture with 5 conv and 2 fc layers with SGD (Simple Gradient descent) as optimizer. And the learning rate was 0.01. I used sample-wise centre as one of the pre-processing techniques. Dropout for fc layers is 0.6. Below are the plots for accuracy and loss. (Green indicates validation. Blue indicates Training). The accuracy after 140 epochs was 75% and 62% respectively for training and validation. We can clearly see that the network is not learning properly and is also overfitting.





- Second. Instead of the traditional SGD optimizer I have experimented with using NADAM optimizer. The performance was worse and the model was not able to learn anything properly so it was stopped at 20 epochs.



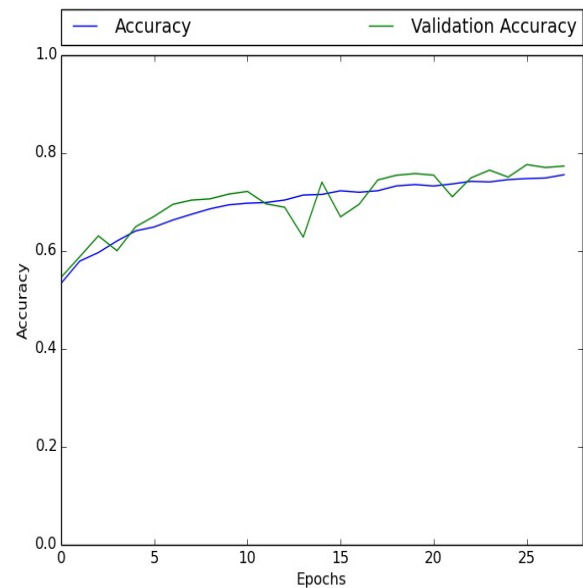

- Third. Then I went back to SGD and increased one convolutional layer and dropout to 0.7 . I decreased the learning rate manually at certain intervals with increase in epochs. There was not much difference between this and the first model.





- Fourth. I felt something was fishy with the sample-wise centering of data and removed it. I used the same architecture as first model but decreased the dropout of fc layers to 0.4 and changed the optimizer to Adadelata. A vast improvement in accuracy and loss is observed.



- Fifth. I confirmed by doubt on the sample-wise centering of input data by using the same model as fourth with processing the input by sample-wise centering. We see that from the below figures that using sample-wise centering of image reduces the performance of our model. Generally sample-wise centering is used to make the input data robust from brightness and contrast changes. But for this dataset, I think this plays an important feature for recognising cats and dogs.



- Sixth. In this model, I added one more convoutional layer to the fourth model. It can be observed that the validation accuracy does not deviate away from training accuracy at higher epochs when compared to 4<sup>th</sup> model due to more denser architecture.



### Other Inferences:

- It can be observed that there are large variations in the validation accuracy due to smaller validation split.
- Adadelta was better in training the model when compared to SGD which can be observed from comparing 4<sup>th</sup> and 3<sup>rd</sup> models.
- Final Training Accuracy: **89.35%** , Final Validation Accuracy: **90.46%**
- Final Loss: **0.251**, Final Validation Loss: **0.249**